

Assignment 3

ST340 Programming for Data Science

Released: Friday week 8, 2017-11-24; Deadline: 11am on Friday week 10, 2017-12-08.

Instructions

- Work in groups of at least ONE and at most TWO.
- Specify your student numbers and names on your assignment.
- Any programming should be in R. Your report should be created using R Markdown and include any code you have written.
- Hand in the compiled assignment, clearly marked with “ST340 Assignment 3” to the Statistics department undergraduate office.

Student number(s):

Student name(s):

Q1

Here is a function that does gradient descent to find local minima:

```
gradient.descent <- function(f, df, x0, iterations=1000, alpha=0.2) {  
  x<-x0  
  for (i in 1:iterations) {  
    cat(i,"/",iterations," : ",x," ",f(x),"\n")  
    x<-x-alpha*df(x)  
  }  
  return(x)  
}
```

Example:

```
f <-function(x) { sum(x^2) }  
df<-function(x) { 2*x }  
gradient.descent(f,df,c(10,20),10,0.2)
```

Q1a

Write a *short* function that uses `gradient.descent` to find a local *maxima*. (For the purpose of this question, `gradient.descent` is a “black box”. Don’t worry about the printed output, just the return value matters.)

i.e.

```
gradient.ascent <- function(f, df, x0, iterations=1000, alpha=0.2) {  
  ... use gradient.descent(...) here ...  
}  
f <-function(x) { (1+x^2)^(-1) }  
df<-function(x) { -2*x*(1+x^2)^(-2) }  
gradient.ascent(f,df,3,40,0.5)
```

Q1b

Consider the function $f : R^2 \rightarrow R$

```
f <- function(x) (x[1]-1)^2 + 100*(x[1]^2-x[2])^2
```

- (1) Give a short mathematical proof that f has a unique minima.
- (2) Write a function df to calculate the gradient of f, i.e.

```
df <- function(x) { ... use x[1] and x[2] ... }
```
- (3) Starting from the point $x_0=c(3,4)$, try to find the minimum using gradient descent.

```
gradient.descent(f,df,c(3,4), ... , ...)
```

Q1c

Write a function to do gradient descent with momentum. Starting from the point $x_0=c(3,4)$, use your function to the minimum of the function from part b.

Q2

Load the tiny MNIST dataset:

```
load("mnist.tiny.RData")
train.X=train.X/255
test.X=test.X/255
```

show some digits:

```
library(grid)
grid.raster( array(aperm(array(train.X[1:50,],c(5,10,28,28)),c(4,1,3,2)),c(140,280)),
             interpolate=F)
```

Use 3-fold cross validation on the training set to compare SVMs with linear kernels, polynomial kernels and RBF kernels. i.e.

```
library(e1071)
svm(train.X,train.labels,type="C-classification",kernel="linear",cross=3)$tot.accuracy
svm(train.X,train.labels,type="C-classification",kernel="poly",
     degree=2,coef=1,cross=3)$tot.accuracy
```

etc, etc.

For the RBF kernels, write a grid search function that takes two lists, `log.C.range` and `log.gamma.range`, and for each pair (lc,lg) of entries in the pair of lists attempts cross-validation with parameters `cost = exp(lc)` and `gamma=exp(lg)`. Once you have found the model with the best cross-validation error, train it on the full training set and then test it on the test set.